



RAJALAKSHMI ENGINEERING COLLEGE

Approved by AICTE | Affiliated to Anna University | Accredited by NAAC

Department of Computer Science and Engineering

CS23334 Fundamentals of Data Science Lab
III semester II Year (2023R)

Name of the Student : Yashwanth Kumar V

Register Number : 240701609



[]

#day-1
#yashwanth kumar v
#240701609
#Matplotlib
#Line plot

[1] ✓ Os

import matplotlib.pyplot as plt
overs=list(range(5,51,5))
India_Score=[400,360,270,365,276,298,350,245,345,299]
Srilanga_Score=[394,367,263,300,277,290,355,202,304,243]
plt.plot(overs,India_Score,'color'=='blue')
plt.plot(overs,Srilanga_Score)
plt.show()
plt.title("India Vs Srilanga")
plt.xlabel("overs")
plt.ylabel("score")
plt.legend()
plt.plot(overs,India_Score,color="blue",label="India")
plt.plot(overs,Srilanga_Score,color="green",label="Srilanga")

/tmp/ipython-input-1389799779.py:11: UserWarning: No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend
[<matplotlib.lines.Line2D at 0x7eda6d5207d0>]

[2] ●

#Line plot
import matplotlib.pyplot as plt
x = [1, 2, 3, 4, 5]
y = [2, 4, 1, 8, 7]
plt.figure(figsize=(6, 4))
plt.plot(x, y, color='blue', marker='o', linestyle='--', label='Line Plot')
plt.title("Line Plot Example")
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.legend()
plt.grid(True)
plt.show()

[3] ✓ Os

BAR CHART
categories = ['Apple', 'Banana', 'Cherry', 'Date']
values = [23, 17, 35, 29]

plt.figure(figsize=(6, 4))
plt.bar(categories, values, color='green')
plt.title("Bar Chart Example")
plt.xlabel("Fruits")
plt.ylabel("Quantity")
plt.show()

[4] ✓ Os

SCATTER PLOT
x_scatter = [5, 7, 8, 7, 2, 17, 2, 9]
y_scatter = [99, 86, 87, 88, 100, 86, 103, 87]
plt.figure(figsize=(6, 4))
plt.scatter(x_scatter, y_scatter, color='red')
plt.title("Scatter Plot Example")
plt.xlabel("X-values")
plt.ylabel("Y-values")
plt.show()

[5] ✓ Os

HISTOGRAM
data = [22, 87, 5, 43, 56, 73, 55, 54, 11, 20, 51, 5, 79, 31, 27]
plt.figure(figsize=(6, 4))
plt.hist(data, bins=5, color='purple', edgecolor='black')
plt.title("Histogram Example")
plt.xlabel("Value Range")
plt.ylabel("Frequency")
plt.show()

[6] ●

PIE CHART
labels = ['Python', 'Java', 'C++', 'Ruby']
sizes = [215, 130, 245, 210]
colors = ['gold', 'lightcoral', 'lightskyblue', 'lightgreen']
explode = (0.1, 0, 0, 0)
plt.figure(figsize=(6, 6))
plt.pie(sizes, explode=explode, labels=labels, colors=colors,autopct='%1.1f%%', shadow=True, startangle=140)
plt.title("Pie Chart Example")
plt.axis('equal')
plt.show()



[]#Exp No:2 Upload and Analyze the data set given in csv format and perform data preprocessing and visualization.
#yashwanth kumar v
#240701609
#experiment-2

[5] 0simport pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

file_path = '/content/sales_data.csv'
df = pd.read_csv(file_path)

print(df.head())
print(df.isnull().sum())

df['Sales'].fillna(df['Sales'].mean(), inplace=True)
df.dropna(subset=['Product', 'Quantity', 'Region'], inplace=True)

print(df.describe())

product_summary = df.groupby('Product').agg({
 'Sales': 'sum',
 'Quantity': 'sum'
}).reset_index()
print(product_summary)

plt.figure(figsize=(10, 6))
plt.bar(product_summary['Product'], product_summary['Sales'])
plt.xlabel('Product')
plt.ylabel('Total Sales')
plt.title('Total Sales by Product')
plt.show()

df['Date'] = pd.to_datetime(df['Date'], format='%d-%m-%Y', errors='coerce')

sales_over_time = df.groupby('Date').agg({'Sales': 'sum'}).reset_index()

plt.figure(figsize=(10, 6))
plt.plot(sales_over_time['Date'], sales_over_time['Sales'])
plt.xlabel('Date')
plt.ylabel('Total Sales')
plt.title('Sales Over Time')
plt.show()

pivot_table = df.pivot_table(values='Sales', index='Region', columns='Product',
 aggfunc=np.sum, fill_value=0)

print(pivot_table)

correlation_matrix = df.corr(numeric_only=True)
print(correlation_matrix)

plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

...
 Date Product Sales Quantity Region
0 01-01-2023 Product A 200 4 North
1 02-01-2023 Product B 150 3 South
2 03-01-2023 Product A 220 5 North
3 04-01-2023 Product C 300 6 East
4 05-01-2023 Product B 180 4 West
Date 0
Product 0
Sales 0
Quantity 0
Region 0
dtype: int64

 Sales Quantity
count 16.000000 16.000000
mean 237.500000 5.375000
std 64.031242 1.746425
min 150.000000 3.000000
25% 187.500000 4.000000
50% 225.000000 5.500000
75% 302.500000 7.000000
max 340.000000 8.000000

 Product Sales Quantity
0 Product A 1350 33
1 Product B 850 17
2 Product C 1600 36
/tmp/ipython-input-2753605132.py:12: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the op

df['Sales'].fillna(df['Sales'].mean(), inplace=True)

Total Sales by Product

Product	Total Sales
Product A	1350
Product B	850
Product C	1600



Q Commands + Code + Text ▶ Run all ↺

RAM
Disk

👤 ⚙️

[]

#Experiment(Handling Missing and Inappropriate Data in a Dataset)
#yashwanth kumar v
#240701069

[1] ✓ Os

import numpy as np
import pandas as pd
df=pd.read_csv("/content/Hotel_Dataset.csv")
df

✓

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

Next steps: [Generate code with df](#) [New interactive sheet](#)

[2] ✓ Os

df.duplicated()

✓

```
0
0 False
1 False
2 False
3 False
4 False
5 False
6 False
7 False
8 False
9 True
10 False

dtype: bool
```

[3] ✓ Os

df.info()

✓ Os

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11 entries, 0 to 10
Data columns (total 9 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   CustomerID          11 non-null    int64
 1   Age_Group           11 non-null    object
 2   Rating(1-5)         11 non-null    int64
 3   Hotel               11 non-null    object
 4   FoodPreference      11 non-null    object
 5   Bill                11 non-null    int64
 6   NoOfPax             11 non-null    int64
 7   EstimatedSalary     11 non-null    int64
 8   Age_Group.1         11 non-null    object
dtypes: int64(5), object(4)
memory usage: 924.0+ bytes
```

[4] ✓ Os

df.drop_duplicates(inplace=True)
df

✓

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

Next steps: [Generate code with df](#) [New interactive sheet](#)

[5] ✓ Os

len(df)

✓ Os

10

[6] ✓ Os

index=np.array(list(range(0,len(df))))
df.set_index(index,inplace=True)
index

✓

array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])

[7] ✓ Os

df

✓

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

Next steps: [Generate code with df](#) [New interactive sheet](#)

[8] ✓ Os

df.drop(['Age_Group.1'],axis=1,inplace=True)
df

✓

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1	20-25	4	Ibis	veg	1300	2	40000
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000
2	3	25-30	6	RedFox	Veg	1322	2	30000
3	4	20-25	-1	LemonTree	Veg	1234	2	120000
4	5	35+	3	Ibis	Vegetarian	989	2	45000
5	6	35+	3	Ibys	Non-Veg	1909	2	122220
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122
7	8	20-25	7	LemonTree	Veg	2999	-10	345673
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777

Next steps: [Generate code with df](#) [New interactive sheet](#)

[9] ✓ Os

df.CustomerID.loc[df.CustomerID<0]=np.nan
df.Bill.loc[df.Bill<0]=np.nan
df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan
df

✓

/tmp/ipython-input-2080958306.py:1: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!
You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use 'df.loc[row_indexer, "col"] = values' instead, to perform the assignment in a single step and ensure this keeps updating the original 'df'.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df.CustomerID.loc[df.CustomerID<0]=np.nan
/tmp/ipython-input-2080958306.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df.CustomerID.loc[df.CustomerID<0]=np.nan
/tmp/ipython-input-2080958306.py:2: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!
You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use 'df.loc[row_indexer, "col"] = values' instead, to perform the assignment in a single step and ensure this keeps updating the original 'df'.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df.Bill.loc[df.Bill<0]=np.nan
/tmp/ipython-input-2080958306.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df.Bill.loc[df.Bill<0]=np.nan
/tmp/ipython-input-2080958306.py:3: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!
You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use 'df.loc[row_indexer, "col"] = values' instead, to perform the assignment in a single step and ensure this keeps updating the original 'df'.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan
/tmp/ipython-input-2080958306.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df.EstimatedSalary.loc[df.EstimatedSalary<0]=np.nan

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4	Ibis	veg	1300.0	2	40000.0
1	2.0	30-35	5	LemonTree	Non-Veg	2000.0	3	59000.0
2	3.0	25-30	6	RedFox	Veg	1322.0	2	30000.0
3	4.0	20-25	-1	LemonTree	Veg	1234.0	2	120000.0

Next steps: [Generate code with df](#) [New interactive sheet](#)

[10] ✓ Os

df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan
df

✓

/tmp/ipython-input-2129877948.py:1: FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!
You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use 'df.loc[row_indexer, "col"] = values' instead, to perform the assignment in a single step and ensure this keeps updating the original 'df'.

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan
/tmp/ipython-input-2129877948.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy.

df['NoOfPax'].loc[(df['NoOfPax']<1) | (df['NoOfPax']>20)]=np.nan

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4	Ibis	veg	1300.0	2.0	40000.0
1	2.0	30-35	5	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	6	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	-1	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3	Ibis	Vegetarian	989.0	2.0	45000.0
5	6.0	35+	3	Ibys	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4	RedFox	Vegetarian	1000.0	NaN	21122.0
7	8.0	20-25	7	LemonTree	Veg	2999.0	NaN	345673.0
8	9.0	25-30	2	Ibis	Non-Veg	3456.0	3.0	NaN
9	10.0	30-35	5	RedFox	non-Veg	NaN	4.0	87777.0

Next steps: [Generate code with df](#) [New interactive sheet](#)

[11] ✓ Os

df.Age_Group.unique()

✓

array(['20-25', '30-35', '25-30', '35+'], dtype=object)

[12] ✓ Os

df.Hotel.unique()

✓

array(['Ibis', 'LemonTree', 'RedFox', 'Ibys'], dtype=object)

[13] ✓ Os

df.Hotel.replace(['Ibys'],'Ibis',inplace=True)
df.FoodPreference.unique

✓

/tmp/ipython-input-456608217.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation on the DataFrame or Series in place.

df.Hotel.replace(['Ibys'],'Ibis',inplace=True)

```
pandas.core.series.Series.unique
def unique() -> ArrayLike
```

Notes

Returns the unique values as a NumPy array. In case of an extension-array backed Series, a new :class:`api.extensions.ExtensionArray` of that type with just the unique values is returned. This includes

[14] ✓ Os

df.FoodPreference.replace(['Vegetarian'],'veg',inplace=True)
df.FoodPreference.replace(['non-Veg'],'Non-Veg',inplace=True)

[15] ✓ Os

df.EstimatedSalary.fillna(round(df.NoOfPax.mean()),inplace=True)
df.NoOfPax.fillna(round(df.NoOfPax.median()),inplace=True)
df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()), inplace=True)
df.Bill.fillna(round(df.Bill.mean()),inplace=True)
df

✓

/tmp/ipython-input-3711388855.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation on the DataFrame or Series in place.

df.EstimatedSalary.fillna(round(df.NoOfPax.mean()),inplace=True)

/tmp/ipython-input-3711388855.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation on the DataFrame or Series in place.

df.NoOfPax.fillna(round(df.NoOfPax.median()),inplace=True)

/tmp/ipython-input-3711388855.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation on the DataFrame or Series in place.

df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()), inplace=True)

/tmp/ipython-input-3711388855.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation on the DataFrame or Series in place.

df.Bill.fillna(round(df.Bill.mean()),inplace=True)

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4	Ibis	Veg	1300.0	2.0	40000.0
1	2.0	30-35	5	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	6	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	-1	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3	Ibis	Veg	989.0	2.0	45000.0
5	6.0	35+	3	Ibis	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4	RedFox	Veg	1000.0	2.0	21122.0
7	8.0	20-25	7	LemonTree	Veg	2999.0	2.0	345673.0
8	9.0	25-30	2	Ibis	Non-Veg	3456.0	3.0	96755.0
9	10.0	30-35	5	RedFox	Non-Veg	1801.0	4.0	87777.0

Next steps: [Generate code with df](#) [New interactive sheet](#)

1

Start coding or [generate with AI](#).

⬆️ ⬇️ ↻ 🗑️ ⋮

🔗 Variables 📄 Terminal

🔄

✓ 5:49 PM 🐍 Python



```
#experiment-4(Experiment to understand the data preprocessing in Data science)
#yashwanth kumar v
#240701609
```

```
[1]
import numpy as np
import pandas as pd
df=pd.read_csv("/content/pre_process_datasample.csv")
df
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	NaN	Yes
5	France	35.0	58000.0	Yes
6	Spain	NaN	52000.0	No
7	France	48.0	79000.0	Yes
8	Germany	50.0	83000.0	No
9	France	37.0	67000.0	Yes

Next steps:

[Generate code with df](#)[New interactive sheet](#)

```
[2]
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Country     10 non-null     object
1   Age         9 non-null      float64
2   Salary      9 non-null      float64
3   Purchased   10 non-null     object
dtypes: float64(2), object(2)
memory usage: 452.0+ bytes
```

```
[3]
df.Country.mode()
```

	Country
0	France

dtype: object

```
[4]
df.Country.mode()[0]
```

'France'

```
[5]
df.Country.fillna(df.Country.mode()[0],inplace=True)
df.Age.fillna(df.Age.median(),inplace=True)
df.Salary.fillna(round(df.Salary.mean()),inplace=True)
df
```

/tmp/ipython-input-1020198583.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.Country.fillna(df.Country.mode()[0],inplace=True)
```

/tmp/ipython-input-1020198583.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.Age.fillna(df.Age.median(),inplace=True)
```

/tmp/ipython-input-1020198583.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df.Salary.fillna(round(df.Salary.mean()),inplace=True)
```

	Country	Age	Salary	Purchased
0	France	44.0	72000.0	No
1	Spain	27.0	48000.0	Yes
2	Germany	30.0	54000.0	No
3	Spain	38.0	61000.0	No
4	Germany	40.0	63778.0	Yes
5	France	35.0	58000.0	Yes
6	Spain	38.0	52000.0	No
7	France	48.0	79000.0	Yes
8	Germany	50.0	83000.0	No
9	France	37.0	67000.0	Yes

Next steps:

[Generate code with df](#)[New interactive sheet](#)

```
[6]
pd.get_dummies(df.Country)
```

	France	Germany	Spain
0	True	False	False
1	False	False	True
2	False	True	False
3	False	False	True
4	False	True	False
5	True	False	False
6	False	False	True
7	True	False	False
8	False	True	False
9	True	False	False

```
[7]
updated_dataset=pd.concat([pd.get_dummies(df.Country),df.iloc[:,[1,2,3]]],axis=1)
updated_dataset
```

	France	Germany	Spain	Age	Salary	Purchased
0	True	False	False	44.0	72000.0	No
1	False	False	True	27.0	48000.0	Yes
2	False	True	False	30.0	54000.0	No
3	False	False	True	38.0	61000.0	No
4	False	True	False	40.0	63778.0	Yes
5	True	False	False	35.0	58000.0	Yes
6	False	False	True	38.0	52000.0	No
7	True	False	False	48.0	79000.0	Yes
8	False	True	False	50.0	83000.0	No
9	True	False	False	37.0	67000.0	Yes

Next steps:

[Generate code with updated_dataset](#)[New interactive sheet](#)

```
[8]
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Country     10 non-null     object
1   Age         10 non-null     float64
2   Salary      10 non-null     float64
3   Purchased   10 non-null     object
dtypes: float64(2), object(2)
memory usage: 452.0+ bytes
```

```
[10]
updated_dataset.Purchased.replace(['No','Yes'],[0,1],inplace=True)
updated_dataset
```

	France	Germany	Spain	Age	Salary	Purchased
0	True	False	False	44.0	72000.0	0
1	False	False	True	27.0	48000.0	1
2	False	True	False	30.0	54000.0	0
3	False	False	True	38.0	61000.0	0
4	False	True	False	40.0	63778.0	1
5	True	False	False	35.0	58000.0	1
6	False	False	True	38.0	52000.0	0
7	True	False	False	48.0	79000.0	1
8	False	True	False	50.0	83000.0	0
9	True	False	False	37.0	67000.0	1

Next steps:

[Generate code with updated_dataset](#)[New interactive sheet](#)

Start coding or [generate](#) with AI.



```
[ ] ▶ #Experiment-5(to detect outliers in a given data set)
#yashwanth kumar v
#240701609

[1] import numpy as np
array=np.random.randint(1,100,16)
array
array[[32, 93, 10, 75, 19, 18, 88, 93, 27, 80, 23, 17, 63, 67, 12, 30]]

[2] array.mean()
np.float64(46.6875)


[3] np.percentile(array,25)
np.float64(18.75)

[4] np.percentile(array,50)
np.float64(31.0)

[5] np.percentile(array,75)
np.float64(76.25)

[6] np.percentile(array,100)
np.float64(93.0)


[8] def outDetection(array):
    sorted(array)
    Q1,Q3=np.percentile(array,[25,75])
    IQR=Q3-Q1
    lr=Q1-(1.5*IQR)
    ur=Q3+(1.5*IQR)
    return lr,ur
lr,ur=outDetection(array)
lr,ur
(np.float64(-67.5), np.float64(162.5))

[9] import seaborn as sns
%matplotlib inline
sns.displot(array)
<seaborn.axisgrid.FacetGrid at 0x786fee4a34d0>



[10] sns.distplot(array)
/tmp/ipython-input-1133588802.py:1: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751

sns.distplot(array)
<Axes: ylabel='Density'>


[11] new_array=array[(array>lr) & (array<ur)]
new_array
array[[32, 93, 10, 75, 19, 18, 88, 93, 27, 80, 23, 17, 63, 67, 12, 30]]

[12] sns.displot(new_array)
<seaborn.axisgrid.FacetGrid at 0x786fb9f63d70>



[13] lr1,ur1=outDetection(new_array)
lr1,ur1
(np.float64(-67.5), np.float64(162.5))

[14] final_array=new_array[(new_array>lr1) & (new_array<ur1)]
final_array
array[[32, 93, 10, 75, 19, 18, 88, 93, 27, 80, 23, 17, 63, 67, 12, 30]]

[15] sns.distplot(final_array)
/tmp/ipython-input-209491988.py:1: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with
similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751

sns.distplot(final_array)
<Axes: ylabel='Density'>

```

[] Start coding or [generate](#) with AI.



[1]
✓ Os

#Experiment-6(Experiment to understand feature scaling)
#yashwanth kumar v
#240701609

[10]
✓ Os

import numpy as np
import pandas as pd
df=pd.read_csv("/content/pre_process_datasample.csv")
df

CountryAgeSalaryPurchased

0France44.072000.0No

1Spain27.048000.0Yes

2Germany30.054000.0No

3Spain38.061000.0No

4Germany40.0NaNYes

5France35.058000.0Yes

6SpainNaN52000.0No

7France48.079000.0Yes

8Germany50.083000.0No

9France37.067000.0Yes

Next steps: [Generate code with df](#) [New interactive sheet](#)

[11]
✓ Os

df.head()

CountryAgeSalaryPurchased

0France44.072000.0No

1Spain27.048000.0Yes

2Germany30.054000.0No

3Spain38.061000.0No

4Germany40.0NaNYes

Next steps: [Generate code with df](#) [New interactive sheet](#)

[12]
✓ Os

df.Country.fillna(df.Country.mode()[0],inplace=True)
features=df.iloc[:, :-1].values

/tmp/ipython-input-3424832005.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df.Country.fillna(df.Country.mode()[0],inplace=True)

[13]
✓ Os

label=df.iloc[:, -1].values

[15]
✓ 4s

from sklearn.impute import SimpleImputer
age=SimpleImputer(strategy="mean",missing_values=np.nan)
Salary=SimpleImputer(strategy="mean",missing_values=np.nan)
age.fit(features[:, [1]])

SimpleImputer ⓘ ?

SimpleImputer()

[16]
✓ Os

Salary.fit(features[:, [2]])

SimpleImputer ⓘ ?

SimpleImputer()

[17]
✓ Os

SimpleImputer()

SimpleImputer ⓘ ?

SimpleImputer()

[18]
✓ Os

features[:, [1]]=age.transform(features[:, [1]])
features[:, [2]]=Salary.transform(features[:, [2]])
features

array([['France', 44.0, 72000.0],
['Spain', 27.0, 48000.0],
['Germany', 30.0, 54000.0],
['Spain', 38.0, 61000.0],
['Germany', 40.0, 63777.77777777778],
['France', 35.0, 58000.0],
['Spain', 38.77777777777778, 52000.0],
['France', 48.0, 79000.0],
['Germany', 50.0, 83000.0],
['France', 37.0, 67000.0]], dtype=object)

[19]
✓ Os

from sklearn.preprocessing import OneHotEncoder
oh = OneHotEncoder(sparse_output=False)
Country=oh.fit_transform(features[:, [0]])
Country

array([[1., 0., 0.],
[0., 0., 1.],
[0., 1., 0.],
[0., 0., 1.],
[0., 1., 0.],
[1., 0., 0.],
[0., 0., 1.],
[1., 0., 0.],
[0., 1., 0.],
[1., 0., 0.]])

[20]
✓ Os

final_set=np.concatenate((Country, features[:, [1, 2]]),axis=1)
final_set

array([[1.0, 0.0, 0.0, 44.0, 72000.0],
[0.0, 0.0, 1.0, 27.0, 48000.0],
[0.0, 1.0, 0.0, 30.0, 54000.0],
[0.0, 0.0, 1.0, 38.0, 61000.0],
[0.0, 1.0, 0.0, 40.0, 63777.77777777778],
[1.0, 0.0, 0.0, 35.0, 58000.0],
[0.0, 0.0, 1.0, 38.77777777777778, 52000.0],
[1.0, 0.0, 0.0, 48.0, 79000.0],
[0.0, 1.0, 0.0, 50.0, 83000.0],
[1.0, 0.0, 0.0, 37.0, 67000.0]], dtype=object)

[21]
✓ Os

from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
sc.fit(final_set)
feat_standard_scaler=sc.transform(final_set)
feat_standard_scaler

array([[1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
7.58874362e-01, 7.49473254e-01],
[-8.16496581e-01, -6.54653671e-01, 1.52752523e+00,
-1.71150388e+00, -1.43817841e+00],
[-8.16496581e-01, 1.52752523e+00, -6.54653671e-01,
-1.27555478e+00, -8.91265492e-01],
[-8.16496581e-01, -6.54653671e-01, 1.52752523e+00,
-1.13023841e-01, -2.53200424e-01],
[-8.16496581e-01, 1.52752523e+00, -6.54653671e-01,
1.77608893e-01, 6.63219199e-16],
[1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
-5.48972942e-01, -5.26656882e-01],
[-8.16496581e-01, -6.54653671e-01, 1.52752523e+00,
0.00000000e+00, -1.07356980e+00],
[1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
1.34013983e+00, 1.38753832e+00],
[-8.16496581e-01, 1.52752523e+00, -6.54653671e-01,
1.63077256e+00, 1.75214693e+00],
[1.22474487e+00, -6.54653671e-01, -6.54653671e-01,
-2.58340208e-01, 2.93712492e-01]])

[22]
✓ Os

from sklearn.preprocessing import MinMaxScaler
mms=MinMaxScaler(feature_range=(0,1))
mms.fit(final_set)
feat_minmax_scaler=mms.transform(final_set)
feat_minmax_scaler

array([[1., 0., 0., 0.73913043, 0.68571429],
[0., 0., 1., 0., 0.],
[0., 1., 0., 0.13043478, 0.17142857],
[0., 0., 1., 0.47826087, 0.37142857],
[0., 1., 0., 0.56521739, 0.45079365],
[1., 0., 0., 0.34782609, 0.28571429],
[0., 0., 1., 0.51207729, 0.11428571],
[1., 0., 0., 0.91304348, 0.88571429],
[0., 1., 0., 1., 1.],
[1., 0., 0., 0.43478261, 0.54285714]])

↑ ↓ ↶ ↷ 🗑 ⋮

CommandsCodeTextRun allSaving failed since 10:04 PMRAMDiskPython 3

[1]0#Experiment-7(EDA)
#yashwanth kumar v
#240701099

[2]import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

[3]tips = sns.load_dataset('tips')
tips.head()

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

Next steps: [Generate code with tips](#) [New interactive sheet](#)

[4]sns.displot(tips['total_bill'], kde=True)

[5]sns.displot(tips['total_bill'], kde=False)

[6]sns.jointplot(x=tips['tip'], y=tips['total_bill'])

[8]sns.jointplot(x=tips['tip'], y=tips['total_bill'], kind='reg')

[7]sns.jointplot(x=tips['tip'], y=tips['total_bill'], kind='hex')

[8]sns.pairplot(tips)

[9]tips.time.value_counts()

time	count
Dinner	176
Lunch	68

dtype: int64

[10]sns.pairplot(tips, hue='time')

[11]sns.pairplot(tips, hue='day')

[12]sns.heatmap(tips.corr(numeric_only=True), annot=True)

<Axes: >

[13]sns.boxplot(tips['total_bill'])

<Axes: ylabel='total_bill'>

[14]sns.boxplot(tips['tip'])

<Axes: ylabel='tip'>

[15]sns.countplot(tips['day'])

<Axes: xlabel='count', ylabel='day'>

Automatic saving failed. This file was updated remotely or in another tab. Show diff

10:06 PMPython 3



```
[ ]
#Experiment-8(linear regression)
#yashwanth kumar v
#240701609
```

```
[1]
import numpy as np
import pandas as pd
```

```
[3]
# Create sample salary dataset
data = {
    'YearsExperience': [1.1, 1.3, 1.5, 2.0, 2.2, 2.9, 3.0, 3.2, 3.2, 3.7, 3.9, 4.0, 4.0, 4.1, 4.5, 4.9, 5.1, 5.3, 5.9, 6.0, 6.8, 7.1, 7.9, 8.2, 8.7, 9.0, 9.5, 9.6, 10.3, 10.5],
    'Salary': [37731, 43525, 39891, 56642, 60150, 54445, 64445, 57189, 63218, 55794, 56957, 57081, 61111, 67938, 66029, 83088, 81363, 93940, 91738, 98273, 101302, 113812, 109431, 105582, 116969, 112635, 122391, 121872, 91738, 93940]
}
df = pd.DataFrame(data)
df
```

	YearsExperience	Salary
0	1.1	37731
1	1.3	43525
2	1.5	39891
3	2.0	56642
4	2.2	60150
5	2.9	54445
6	3.0	64445
7	3.2	57189
8	3.2	63218
9	3.7	55794
10	3.9	56957
11	4.0	57081
12	4.0	61111
13	4.1	67938
14	4.5	66029
15	4.9	83088
16	5.1	81363
17	5.3	93940
18	5.9	91738
19	6.0	98273
20	6.8	101302
21	7.1	113812
22	7.9	109431
23	8.2	105582
24	8.7	116969
25	9.0	112635
26	9.5	122391
27	9.6	121872
28	10.3	91738
29	10.5	93940

```
[4]
df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
Column Non-Null Count Dtype
--- ---
0 YearsExperience 30 non-null float64
1 Salary 30 non-null int64
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes

```
[5]
df.dropna(inplace=True)
df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
Column Non-Null Count Dtype
--- ---
0 YearsExperience 30 non-null float64
1 Salary 30 non-null int64
dtypes: float64(1), int64(1)
memory usage: 612.0 bytes

```
[6]
df.describe()
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	79340.666667
std	2.837888	26128.480351
min	1.100000	37731.000000
25%	3.200000	57108.000000
50%	4.700000	74650.500000
75%	7.700000	100544.750000
max	10.500000	122391.000000

```
[7]
features = df.iloc[:, 0].values
label = df.iloc[:, 1].values
```

```
[8]
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=0)
```

```
[9]
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train.reshape(-1, 1), y_train)
```

LinearRegression

LinearRegression()

```
[10]
model.score(x_train.reshape(-1, 1), y_train)
```

0.8308613552280183

```
[11]
model.score(x_test.reshape(-1, 1), y_test)
```

0.7657841556832163

```
[12]
model.coef_
```

array([8575.37154813])

```
[13]
model.intercept_
```

np.float64(35136.91225239668)

```
[14]
import pickle
pickle.dump(model, open('SalaryPred.model', 'wb'))
```

```
[15]
model = pickle.load(open('SalaryPred.model', 'rb'))
yprof_exp = float(input('Enter Years of Experience '))
yprof_expNP = np.array(yprof_exp)
Salary = model.predict(yprof_expNP.reshape(-1, 1))
print('Estimated Salary for {} years of experience is {}'.format(yprof_exp, Salary))
```

Enter Years of Experience 5
Estimated Salary for 5.0 years of experience is [78013.76999303]

The screenshot shows a file explorer interface. The title bar reads "Files". On the left is a sidebar with icons for a menu, search, navigation (back/forward), and file operations (copy, paste, delete, rename). The main pane displays a file tree with the following structure:

- .. (parent directory)
- sample_data (folder)
- Social_Network_Ads.csv (CSV file)

A toolbar at the top of the main pane includes icons for upload, refresh, create folder, and delete.

[]

#Experiment-9(Logistic regression)
#yashwanth kumar v
#240701609

[9]
✓ Os

import numpy as np
import pandas as pd

[10]
✓ Os

df = pd.read_csv('/content/Social_Network_Ads.csv')
df

User IDGenderAgeEstimatedSalaryPurchased

015624510Male19190000

115810944Male35200000

215668575Female26430000

315603246Female27570000

415804002Male19760000

...

39515691863Female46410001

39615706071Male51230001

39715654296Female50200001

39815755018Male36330000

39915594041Female49360001

400 rows × 5 columns

Next steps: [Generate code with df](#) [New interactive sheet](#)

[11]
✓ Os

df.head()

User IDGenderAgeEstimatedSalaryPurchased

015624510Male19190000

115810944Male35200000

215668575Female26430000

315603246Female27570000

415804002Male19760000

Next steps: [Generate code with df](#) [New interactive sheet](#)

[12]
✓ Os

features = df.iloc[:, [2, 3]].values
label = df.iloc[:, 4].values

[13]
✓ Os

features

Show hidden output

[14]
✓ Os

label

array([[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0,
1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1,
0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1,
1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,
0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1,
0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,
1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,
1, 1, 0, 1])

[15]
✓ Ts

from sklearn.model_selection import train_test_split

[16]
✓ Os

from sklearn.linear_model import LogisticRegression

[17]
✓ 2s

Find best random_state
for i in range(1, 40):
x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=i)
model = LogisticRegression()
model.fit(x_train, y_train)
train_score = model.score(x_train, y_train)
test_score = model.score(x_test, y_test)
if test_score > train_score:
print('Test {} Train {} Random State {}'.format(test_score, train_score, i))

Test 0.9 Train 0.840625 Random State 4
Test 0.8625 Train 0.85 Random State 5
Test 0.8625 Train 0.859375 Random State 6
Test 0.8875 Train 0.8375 Random State 7
Test 0.8625 Train 0.8375 Random State 9
Test 0.9 Train 0.840625 Random State 10
Test 0.8625 Train 0.85625 Random State 14
Test 0.85 Train 0.84375 Random State 15
Test 0.8625 Train 0.85625 Random State 16
Test 0.875 Train 0.834375 Random State 18
Test 0.85 Train 0.84375 Random State 19
Test 0.875 Train 0.84375 Random State 20
Test 0.8625 Train 0.834375 Random State 21
Test 0.875 Train 0.840625 Random State 22
Test 0.875 Train 0.840625 Random State 24
Test 0.85 Train 0.834375 Random State 26
Test 0.85 Train 0.840625 Random State 27
Test 0.8625 Train 0.834375 Random State 30
Test 0.8625 Train 0.85625 Random State 31
Test 0.875 Train 0.853125 Random State 32
Test 0.8625 Train 0.84375 Random State 33
Test 0.875 Train 0.83125 Random State 35
Test 0.8625 Train 0.853125 Random State 36
Test 0.8875 Train 0.840625 Random State 38
Test 0.875 Train 0.8375 Random State 39

[18]
✓ Os

Train final model with best random_state
x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=25)
final_model = LogisticRegression()
final_model.fit(x_train, y_train)
print(final_model.score(x_train, y_train))
print(final_model.score(x_test, y_test))

0.846875
0.8

[19]
✓ Os

from sklearn.metrics import classification_report
print(classification_report(label, final_model.predict(features)))

precisionrecallf1-score support

00.840.920.88257
10.820.690.75143

accuracy0.84400
macro avg0.830.810.82400
weighted avg0.840.840.83400

[]

Start coding or generate with AI.

Files

sample_data

Mall_Customers.csv

#Experiment-10(K-Means Clustering)
#yashwanth kumar v
#240701609

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

df = pd.read_csv('/content/Mall_Customers.csv')
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- -
0 CustomerID 200 non-null int64
1 Gender 200 non-null object
2 Age 200 non-null int64
3 Annual Income (k\$) 200 non-null int64
4 Spending Score (1-100) 200 non-null int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB

df.head()

CustomerID Gender Age Annual Income (k\$) Spending Score (1-100)
0 1 Male 19 15 39
1 2 Male 21 15 81
2 3 Female 20 16 6
3 4 Female 23 16 77
4 5 Female 31 17 40

Next steps: Generate code with df New interactive sheet

sns.pairplot(df)

<seaborn.axisgrid.PairGrid at 0x7987511f63f0>

features = df.iloc[:, [3, 4]].values

from sklearn.cluster import KMeans
model = KMeans(n_clusters=5)
model.fit(features)

KMeans
KMeans(n_clusters=5)

Final = df.iloc[:, [3, 4]]
Final['label'] = model.predict(features)
Final.head()

/tmp/ipython-input-1374593914.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
Final['label'] = model.predict(features)
Final

Annual Income (k\$) Spending Score (1-100) label
0 15 39 0
1 15 81 3
2 16 6 0
3 16 77 3
4 17 40 0

Next steps: Generate code with Final New interactive sheet

sns.set_style('whitegrid')
sns.FacetGrid(Final, hue='label', height=8).map(plt.scatter, 'Annual Income (k\$)', 'Spending Score (1-100)').add_legend()
plt.show()

features_el = df.iloc[:, [2, 3, 4]].values
from sklearn.cluster import KMeans
wcss = []
for i in range(1, 10):
model = KMeans(n_clusters=i)
model.fit(features_el)
wcss.append(model.inertia_)
plt.plot(range(1, 10), wcss)

[<matplotlib.lines.Line2D at 0x7987491e28a0>]

Start coding or generate with AI.

Files

📁 ..

📁 sample_data

📄 Iris (1).csv

[]

🔍

#Experiment-11(KNN)
#yashwanth kumar v
#240701609

[29]

✓ Os

import numpy as np
import pandas as pd

[30]

✓ Os

df = pd.read_csv('/content/Iris (1).csv')
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
Column Non-Null Count Dtype

0 sepal.length 150 non-null float64
1 sepal.width 150 non-null float64
2 petal.length 150 non-null float64
3 petal.width 150 non-null float64
4 variety 150 non-null object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB

[31]

✓ Os

df.variety.value_counts()

count
variety
Setosa 50
Versicolor 50
Virginica 50
dtype: int64

[32]

✓ Os

df.head()

sepal.length sepal.width petal.length petal.width variety
0 5.1 3.5 1.4 0.2 Setosa
1 4.9 3.0 1.4 0.2 Setosa
2 4.7 3.2 1.3 0.2 Setosa
3 4.6 3.1 1.5 0.2 Setosa
4 5.0 3.6 1.4 0.2 Setosa

Next steps: [Generate code with df](#) [New interactive sheet](#)

[33]

✓ Os

features = df.iloc[:, :-1].values
label = df.iloc[:, 4].values

[34]

✓ Os

from sklearn.model_selection import train_test_split

[35]

✓ Os

from sklearn.neighbors import KNeighborsClassifier

[36]

✓ Os

x_train, x_test, y_train, y_test = train_test_split(features, label, test_size=0.2, random_state=25)
modelKNN = KNeighborsClassifier(n_neighbors=5)
modelKNN.fit(x_train, y_train)

▾ KNeighborsClassifier ⓘ ?

KNeighborsClassifier()

[37]

✓ Os

print(modelKNN.score(x_train, y_train))
print(modelKNN.score(x_test, y_test))

0.9833333333333333
0.9333333333333333

[38]

✓ Os

from sklearn.metrics import confusion_matrix
confusion_matrix(label, modelKNN.predict(features))

array([[50, 0, 0],
[0, 47, 3],
[0, 1, 49]])

[39]

✓ Os

from sklearn.metrics import classification_report
print(classification_report(label, modelKNN.predict(features)))

precision recall f1-score support
Setosa 1.00 1.00 1.00 50
Versicolor 0.98 0.94 0.96 50
Virginica 0.94 0.98 0.96 50
accuracy 0.97 0.97 0.97 150
macro avg 0.97 0.97 0.97 150
weighted avg 0.97 0.97 0.97 150

[]

Start coding or [generate](#) with AI.

📄 Variables 📄 Terminal

69.57 GB available

🌟

✓ 5:31 AM 📄 Python 3

Files

📁 ..

📁 sample_data

📄 Iris (1).csv

[]

#Yashwanth kumar v
#240701609

[40]

Experiment-12(T-TEST)
Testing whether average marks = 70 at 5% significance level

import numpy as np
from scipy import stats

Sample data
marks = np.array([72, 68, 75, 70, 74, 69, 71, 73, 70, 72])

Hypothesized mean
mu0 = 70

One-sample t-test
t_stat, p_value = stats.ttest_1samp(marks, mu0)

print(f"T-statistic: {t_stat:.3f}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
 print("Reject Null Hypothesis: Mean is significantly different from 70.")
else:
 print("Fail to Reject Null Hypothesis: No significant difference.")

▼

T-statistic: 1.993
P-value: 0.0774
Fail to Reject Null Hypothesis: No significant difference.

[41]

#Experiment -13 Z-TEST
Testing manufacturer's claim that average weight = 50g

import numpy as np
from math import sqrt
from scipy.stats import norm

Given data
x_bar = 51.2 # sample mean
mu0 = 50 # population mean
sigma = 3 # population standard deviation
n = 36 # sample size

Calculate Z-statistic
z_stat = (x_bar - mu0) / (sigma / sqrt(n))

Two-tailed p-value
p_value = 2 * (1 - norm.cdf(abs(z_stat)))

print(f"Z-statistic: {z_stat:.3f}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
 print("Reject Null Hypothesis: Mean is significantly different from 50 g.")
else:
 print("Fail to Reject Null Hypothesis: No significant difference.")

▼

... Z-statistic: 2.400
P-value: 0.0164
Reject Null Hypothesis: Mean is significantly different from 50 g.

[42]

Experiment-14 ANOVA TEST
Testing if there's significant difference among three fertilizers (A, B, C)

import numpy as np
from scipy import stats

Data for three fertilizers
A = [20, 22, 23]
B = [19, 20, 18]
C = [25, 27, 26]

Perform one-way ANOVA
f_stat, p_value = stats.f_oneway(A, B, C)

print(f"F-statistic: {f_stat:.3f}")
print(f"P-value: {p_value:.4f}")

alpha = 0.05
if p_value < alpha:
 print("Reject Null Hypothesis: Means are significantly different.")
else:
 print("Fail to Reject Null Hypothesis: No significant difference.")

▼

... F-statistic: 25.923
P-value: 0.0011
Reject Null Hypothesis: Means are significantly different.

[]

Start coding or generate with AI.

Disk69.57 GB available

Automatic saving failed. This file was updated remotely or in another tab. Show diff

✓ 5:36 AM

Python 3